

Assignment-2

February 6, 2025

CVSN Assignment-2

Mayank Nagar

MDS202334

```
[1]: import cv2 as cv
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from collections import Counter, defaultdict
```

1 Question-1

```
[2]: img1 = cv.imread('./images/image1.jpg')
img2 = cv.imread('./images/image2.jpg')

assert img1 is not None
assert img2 is not None
```

```
[3]: plt.figure(figsize=(10,5))
plt.subplot(1,2,1)
plt.imshow(img1[:,:,:-1])
plt.title('Image 1')
plt.axis('off')


plt.subplot(1,2,2)
plt.imshow(img2[:,:,:-1])
plt.title('Image 2')
plt.axis('off')

plt.show()
```




1.0.1 (a) Histograms of each channel of Image-1

```
[4]: plt.figure(figsize=(15,4))
for i, channel in enumerate(['B', 'G', 'R']):
    plt.subplot(1,3,i+1)
    plt.title(f'Histogram of {channel} Channel')
    plt.hist(img1[:, :, i].ravel(), bins=256, fc='k', ec='k')
plt.show()
```



1.0.2 Histograms of each channel of Image-2

```
[5]: plt.figure(figsize=(15,4))
for i, channel in enumerate(['B', 'G', 'R']):
    plt.subplot(1,3,i+1)
    plt.title(f'Histogram of {channel} Channel')
    plt.hist(img2[:, :, i].ravel(), bins=256, fc='k', ec='k')
plt.show()
```



1.0.3 (b) Histogram Matching

```
[6]: def histogram_matching(source, reference):
    # Placeholder
    transformed_img = np.zeros_like(source, dtype=np.uint8)

    # Equalize both images
    he1 = cv.equalizeHist(source)
    he2 = cv.equalizeHist(reference)

    # Transformation Functions
    Tr = {r: tr for r, tr in zip(source.ravel(), he1.ravel())}
    Gz = {z: gz for z, gz in zip(reference.ravel(), he2.ravel())}
    Gz_inv = defaultdict(set)
    for z, gz in Gz.items():
        Gz_inv[gz].add(z)

    # Histogram matching function
    mapping = {}
    for i in range(256):
        v = None
        if i in Tr:
            v = Tr[i]
        else:
            continue

        if v in Gz_inv:
            mapping[i] = list(Gz_inv[v])[0]
        else:
            l = 0
            while v - l >= 0 or v + l <= 255:
                if v - l in Gz_inv and v + l in Gz_inv:
                    mapping[i] = min(Gz_inv[v-l]) # I choose to go with min
                    ↪ value
```

```

        break
    elif v - 1 in Gz_inv:
        mapping[i] = min(Gz_inv[v-1])
        break
    elif v + 1 in Gz_inv:
        mapping[i] = min(Gz_inv[v+1])
        break
    else:
        l += 1

# Matching the source image to reference
for i in range(source.shape[0]):
    for j in range(source.shape[1]):
        transformed_img[i, j] = mapping[source[i, j]]

return mapping, transformed_img

```

```

[7]: # Applying the histogram matching function
transformed_img = np.zeros_like(img1, dtype=np.uint8)
for channel in range(3):
    mapping, matched_channel = histogram_matching(img1[:, :, channel], img2[:, :,
↪ channel])
    transformed_img[:, :, channel] = matched_channel

```

```

[8]: plt.figure(figsize=(15,8))
plt.subplot(1,3,1)
plt.imshow(img1[:, :, :-1])
plt.title('Source Image')
plt.axis('off')

plt.subplot(1,3,2)
plt.imshow(img2[:, :, :-1])
plt.title('Target Image')
plt.axis('off')

plt.subplot(1,3,3)
plt.imshow(transformed_img[:, :, :-1])
plt.title('Histogram Matched Image of Source Image')
plt.axis('off')

plt.show()

```



The reconstructed image appears greenish because the target image contains a high concentration of green pixels.

2 Question-2


2.0.1 (a) Experimenting with different clip limits in contrast limited AHE

```
[9]: xray = cv.imread('./images/chestxray1.png', 0)
      assert xray is not None

      plt.imshow(xray, cmap='gray')
      plt.axis('off')
      plt.show()
```



```
[10]: plt.figure(figsize= (18, 5))
for i, clipLimit in enumerate(range(2,41,4)):
    clahe = cv.createCLAHE(clipLimit=clipLimit)
    plt.subplot(2, 5, i+1)
    img = clahe.apply(xray)
    plt.imshow(img, cmap='gray')
    plt.title(f"{clipLimit = }")
    plt.axis('off')
```



As we increase clipLimit value contrast increases

2.0.2 (b) Implementation of Sliding Window AHE

```
[11]: def SWAHE(image, k):
    h, w = image.shape
    pad = k // 2
    swahe_image = np.zeros_like(image)

    # Pad the image to handle border effects
    padded_image = np.pad(image, pad, mode='reflect')

    # Extract the k*k window centered at each pixel
    for y in range(h):
        for x in range(w):
            window = padded_image[y:y + k, x:x + k]

            # Equalize window
            equalized_window = cv.equalizeHist(window)


            # Map pixel intensity using the computed CDF
            swahe_image[y, x] = equalized_window[pad, pad]

    return swahe_image
```

```
[12]: # Display the results
plt.figure(figsize=(20, 10))
plt.subplot(1, 5, 1)
plt.imshow(xray, cmap='gray')
plt.title("Original Image")
plt.axis("off")

plt.subplot(1, 5, 2)
clahe = cv.createCLAHE(clipLimit=38)
clahe_image = clahe.apply(xray)
plt.imshow(clahe_image, cmap='gray')
plt.title("CLAHE image")
plt.axis("off")

for i, k in enumerate([6,11,17]):
    plt.subplot(1, 5, i+3)
    swahe_image = SWAHE(xray, k)
    plt.imshow(swahe_image, cmap='gray')
    plt.title(f"SWAHE image with window size {k}")
    plt.axis("off")
plt.show()
```



CLAHE enhances the image contrast, but the center (spine) remains overly bright, making it difficult to see details. This issue is resolved by using the SWAHE technique, where windowed histogram equalization ensures more uniform brightness.

When comparing output images from SWAHE, increasing the `window size` improves clarity, making all artifacts more distinguishable.


```
[13]: wiki = cv.imread('images/Unequalized_Hawkes_Bay_NZ.jpg', 0)

# Display the results
plt.figure(figsize=(20, 10))
plt.subplot(1, 5, 1)
plt.imshow(wiki, cmap='gray')
plt.title("Original Image")
plt.axis("off")

plt.subplot(1, 5, 2)
clahe = cv.createCLAHE(clipLimit=38)
```

```
clahe_image = clahe.apply(wiki)
plt.imshow(clahe_image, cmap='gray')
plt.title("CLAHE image")
plt.axis("off")

for i, k in enumerate([6,11,17]):
    plt.subplot(1, 5, i+3)
    swahe_image = SWAHE(wiki, k)
    plt.imshow(swahe_image, cmap='gray')
    plt.title(f"SWAHE image with window size {k}")
    plt.axis("off")
plt.show()
```



In this example of the CLAHE contrast of the image increases and the image becomes clearer. But in the case of the SWAHE, it introduces noise.

2.0.3 (d) Observations about when each method may be preferable

Feature	CLAHE	SWAHE
Noise Handling	Reduces noise by clipping histogram peaks	Can amplify noise, especially in flat regions
Computation Speed	Faster	Slower (Window-based operations are computationally expensive)
Contrast Adjustment	Controlled contrast (via clipLimit)	Can over-enhance contrast in some areas
Edge Preservation	Better at preserving edges	May introduce artifacts at strong edges

3 Question-3

3.0.1 Function for intensity slicing

```
[14]: def intensity_slicing(img, min_range, max_range, method='zero'):
    output = img.copy()
    if method == 'zero':
        output[(min_range >= img) | (img >= max_range)] = 0
        output[(min_range < img) & (img < max_range)] = 255
    elif method == 'unchanged':
        output[(min_range < img) & (img < max_range)] = 255
```



```


else:
    raise ValueError(f'{method} is not valid')
return output

```

```

[15]: kidney_image = cv.imread('./images/kidney.jpg', 0)
      # Display the grayscale face image
      plt.imshow(kidney_image, cmap='gray')
      plt.title("Scan of Kidney")
      plt.axis("off")
      plt.show()

```



3.0.2 (a) highlight a particular range and set the rest to zero;

3.0.3 (b) highlight a particular range and leave the rest unchanged.

```

[16]: plt.figure(figsize=(10, 10))

      plt.subplot(3,2,1)
      plt.imshow(kidney_image, cmap='gray')
      plt.axis('off')
      plt.title('Original kidney_image')

```

```

plt.subplot(3,2,3)

low = 175
high = 255

new_img = intensity_slicing(kidney_image, low, high, method='zero')
plt.imshow(new_img, cmap='gray')
plt.title('Transformed kidney_image in (a)')
plt.axis('off')


plt.subplot(3,2,4)
plt.plot(range(256), [0]*low+[255]*(high-low+1)+[0]*(255-high))
plt.title('Intensity Transformation in (a)')

plt.subplot(3,2,5)
new_img = intensity_slicing(kidney_image, low, high, method='unchanged')
plt.title('Transformed kidney_image in (b)')
plt.imshow(new_img, cmap='gray')
plt.axis('off')

plt.subplot(3,2,6)
plt.plot(range(256), list(range(low))+[255]*(high-low+1)+list(range(high+1,
↪256)))
plt.title('Intensity Transformation in (b)')
plt.show()

```


Original kidney_image




Transformed kidney_image in (a)




Transformed kidney_image in (b)



Intensity Transformation in (a)



Intensity Transformation in (b)



4 Question-4

4.0.1 (a) Extracting and plotting all 8 bit planes

```
[17]: # Create a list to store the bit planes
bit_planes = []

# Extract each bit plane (0-7)
for i in range(8):
```

```

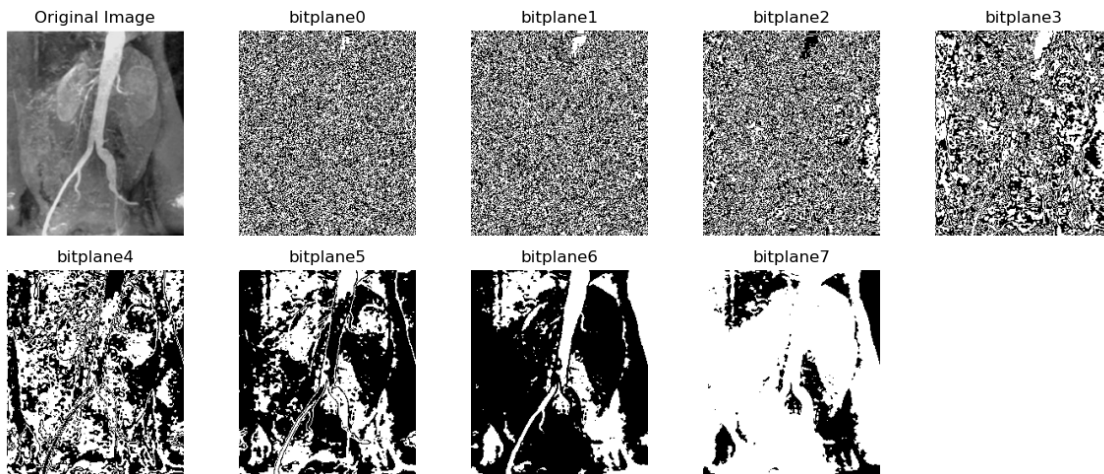
# Use bitwise AND and shift to extract the i-th bit
bit_plane = (kidney_image & (1 << i)) >> i
bit_planes.append(bit_plane)

# Plot the original image and its bit planes
plt.figure(figsize=(12, 5))
plt.subplot(2, 5, 1)
plt.imshow(kidney_image, cmap="gray")
plt.title("Original Image")
plt.axis("off")

# Plotting the bit planes
for i in range(8):
    plt.subplot(2, 5, i + 2)
    plt.imshow(bit_planes[i], cmap="gray")
    plt.title(f"bitplane{i}")
    plt.axis("off")

plt.tight_layout()
plt.show()

```



4.0.2 (b) Plotting intensity transformations used in bitplane0, bitplane3, bitplane7

```

[18]: x_values = np.arange(256)

# Plot the transformations
for bit in [0, 3, 7]:
    plt.figure(figsize=(10, 5))

    # Compute transformations for bitplane 0, 3, and 7

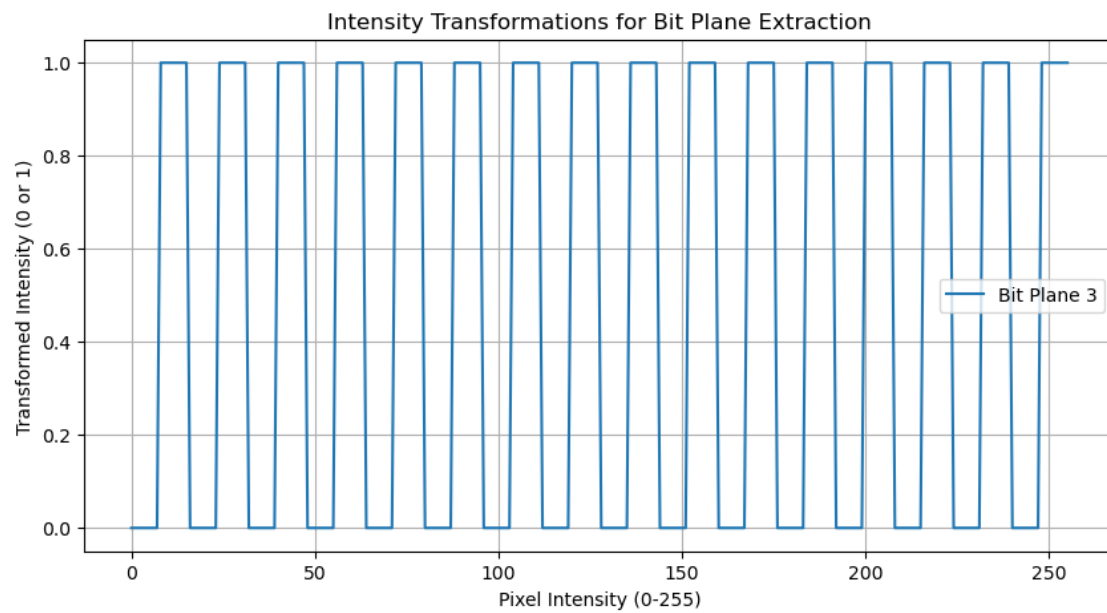

```

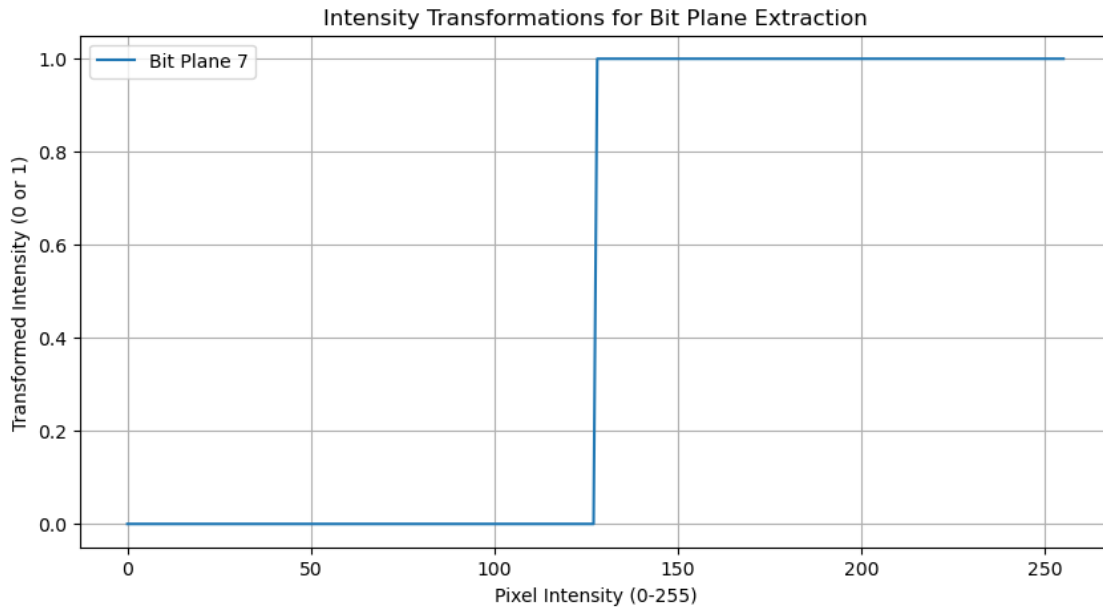
```

y_bitplane = (x_values & (1 << bit)) >> bit

plt.plot(x_values, y_bitplane, label=f"Bit Plane {bit}")
plt.xlabel("Pixel Intensity (0-255)")
plt.ylabel("Transformed Intensity (0 or 1)")
plt.title("Intensity Transformations for Bit Plane Extraction")
plt.legend()
plt.grid(True)
plt.show()

```

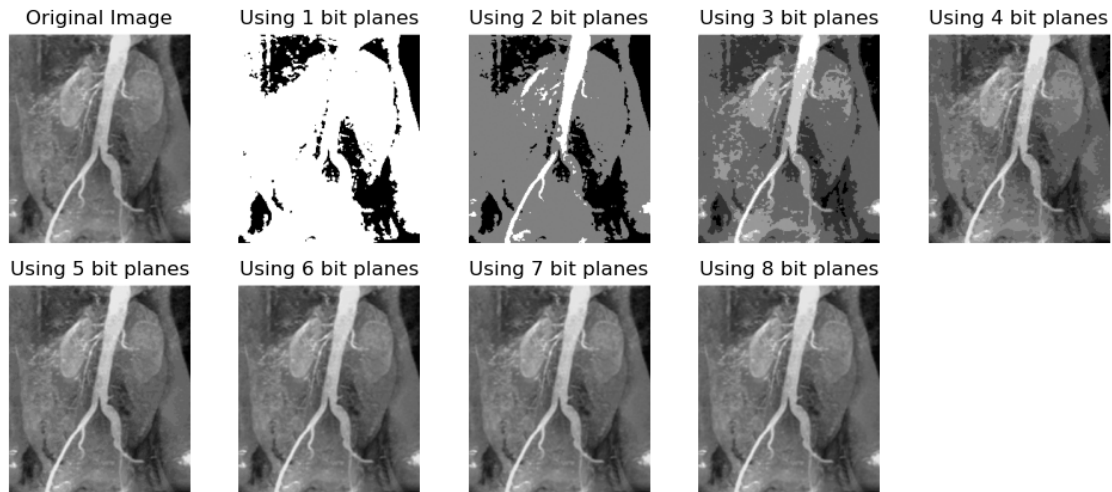




4.0.3 (c) Reconstructing the image with fewer bit planes

```
[19]: plt.figure(figsize=(12, 5))
plt.subplot(2, 5, 1)
plt.imshow(kidney_image, cmap="gray")
plt.title("Original Image")
plt.axis("off")

reconstructed_image = np.zeros_like(kidney_image, dtype=np.uint8)
for i in range(8):
    plt.subplot(2, 5, i+2)
    reconstructed_image = (reconstructed_image*2 + bit_planes[-i-1]).astype(np.
        uint8)
    plt.imshow(reconstructed_image, cmap='gray')
    plt.axis('off')
    plt.title(f'Using {i+1} bit planes')
plt.show()
```



At least 5 bit planes are required to reconstruct the image without losing too much details

5 Question-5

```
[20]: windmill = cv.imread('./images/windmills.jpg', 0)
      assert windmill is not None

      plt.figure(figsize=(5, 15))
      plt.imshow(windmill, cmap='gray')
      plt.title("Windmills - Grayscale Image")
      plt.axis("off")
      plt.show()
```

Windmills - Grayscale Image



5.0.1 (a) Computing the best alignment of two channels of windmill image

```
[21]: def ssd(image1, image2):  
    """Function to compute SSD (Sum of Squared Differences)"""  
    return np.sum((image1 - image2) ** 2)  
  
def best_alignment(fixed, moving, search_range=10):  
    """Function to find the best alignment for a channel"""  
    best_shift = (0, 0)  
    min_ssd = float('inf')  
  
    # Try all displacements in the given range [-search_range, search_range]  
    for dx in range(-search_range, search_range + 1):  
        for dy in range(-search_range, search_range + 1):  
            # Shift the moving image  
            shifted = np.roll(moving, (dy, dx), axis=(0, 1))  
  
            # Compute SSD with the fixed image (blue channel)  
            error = ssd(fixed, shifted)  
  
            # Update if a better alignment is found  
            if error < min_ssd:  
                min_ssd = error  
                best_shift = (dx, dy)  
  
    return best_shift, min_ssd
```

```
[22]: height = windmill.shape[0] // 3  
  
    # Split into three channels (B, G, R)  
    B = windmill[0:height, :]  
    G = windmill[height:2*height, :]  
    R = windmill[2*height:3*height, :]
```

```
[23]: # Align Green and Red channels to the fixed Blue channel  
    (dx_g, dy_g), min_ssd_g = best_alignment(B, G)  
    (dx_r, dy_r), min_ssd_r = best_alignment(B, R)  
  
    # Apply the best shifts  
    aligned_G = np.roll(G, (dy_g, dx_g), axis=(0, 1))  
    aligned_R = np.roll(R, (dy_r, dx_r), axis=(0, 1))
```

5.0.2 (b) Best SSD value and displacement vector for each channel

```
[24]: print(f'Best SSD value for Green channel {min_ssd_g}')
```

```
print(f'Displacement vector for the Green channel {(dx_g, dy_g)}')
```

```
print()
```

```
print(f'Best SSD value for Red channel {min_ssd_r}')
```

```
print(f'Displacement vector for the Red channel {(dx_r, dy_r)}')
```

Best SSD value for Green channel 10118958

Displacement vector for the Green channel (0, -1)

Best SSD value for Red channel 13107332

Displacement vector for the Red channel (0, 10)

5.0.3 (c) Merge aligned channels into a single RGB image

```
[25]: aligned_image = np.dstack([aligned_R, aligned_G, B])
```

5.0.4 (d) Displaying the coloured image

```
[26]: plt.figure(figsize=(10, 5))
```

```
plt.imshow(aligned_image)
```

```
plt.axis("off")
```

```
plt.title("Aligned Image")
```

```
plt.show()
```

Aligned Image



[]: